

# STUDENTS PERFORMANCE PREDICTION

## Introduction

Education is an important element of the society, every government and country in the world work so hard to improve this sector. With the corona-virus outbreak that has disrupted life around the globe in 2020, the educational systems have been affected in many ways; studies show that student's performance has decreased since then, which highlights the need to deal with this problem more seriously and try to find effective solutions, as well as the influencing factors.

## Motivation

The educational systems need, at this specific time, innovative ways to improve quality of education to achieve the best results and decrease the failure rate.

As students of the IT department who studied in the last month a little about machine learning, we know that in order for an institute to provide quality education to learners, deep analysis of previous records of the learners can play a vital role, and wanted to work on this challenging task.

## Objectives

In this notebook, we will:

- **Predict whether or not a student will pass the final exam based on certain information given**
- **Compare the three learning algorithms**
- **Find out what most affects student achievement**
- **Find the best algorithm with high accuracy**

We will be using three learning algorithms:

- **Logistic regression**
- **Supported vector machine**
- **KNN**

## Problem Statement:

As already mentioned, with the help of the old students records, we can come up with a model that can let us help students improve their performance in exams by predicting the student success. So, it is obvious it's a problem of classification, and we will classify a student based on his given informations, and we will also use different classifiers such as KNN or SVM classifier and compare between them. Many factors affect a student performance in exams like family problems or alcohol consumption, and by using our skills in machine learning we want to:

- 1) predict whether a student will pass his final exam or not.
- 2) came up with the best classifier that is more accurate and avoid overfitting and underfitting by using simple techniques.
- 3) know what the most factors affect a student performance.

So, teachers and parents will be able to intervene before students reach the exam stage and solve the problems.

## Dataset:

Dataset name: Student.csv

Source: Kaggle

## EDA:

```
df.shape
```

```
(395, 31)
```

```
df.dropna().shape # their is no null value "fortunately:"
```

```
(395, 31)
```

```
df.columns
```

```
Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',  
      'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',  
      'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',  
      'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',  
      'Walc', 'health', 'absences', 'passed'],  
      dtype='object')
```



```
features=['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
          'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
          'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
          'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
          'Walc', 'health', 'absences']
```

```
#plot of student status
dfv['passed'].value_counts()
```


```
yes      265
no       130
Name: passed, dtype: int64
```

```
df["goout"].unique()
```

```
array([0.75, 0.5 , 0.25, 0.  , 1.  ])
```



## Good Performance vs. Poor Performance Student Weekend Alcohol Consumption



## Pre-processing (dimensionality reduction):

- Scaling:

```
# Let's scal our features
feature_scaling(df)
```

```
# Now we are ready for models training
df
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	absences	passed
0	0.0	1.0	0.059264	0.0	1.0	1.0	1.00	1.00	0.75	0.00	...	0.0	0.0	0.75	0.50	0.75	0.00	0.00	0.50	0.003882	0.0
1	0.0	1.0	0.013809	0.0	1.0	0.0	0.25	0.25	0.75	1.00	...	1.0	0.0	1.00	0.50	0.50	0.00	0.00	0.50	-0.022785	0.0
2	0.0	1.0	-0.077100	0.0	0.0	0.0	0.25	0.25	0.75	1.00	...	1.0	0.0	0.75	0.50	0.25	0.25	0.50	0.50	0.057215	1.0
3	0.0	1.0	-0.077100	0.0	1.0	0.0	1.00	0.50	0.25	0.50	...	1.0	1.0	0.50	0.25	0.25	0.00	0.00	1.00	-0.049451	1.0
4	0.0	1.0	-0.031646	0.0	1.0	0.0	0.75	0.75	1.00	1.00	...	0.0	0.0	0.75	0.50	0.25	0.00	0.25	1.00	-0.022785	1.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
390	1.0	0.0	0.150173	0.0	0.0	1.0	0.50	0.50	0.50	0.50	...	0.0	0.0	1.00	1.00	0.75	0.75	1.00	0.75	0.070549	0.0
391	1.0	0.0	0.013809	0.0	0.0	0.0	0.75	0.25	0.50	0.50	...	1.0	0.0	0.25	0.75	1.00	0.50	0.75	0.25	-0.036118	1.0
392	1.0	0.0	0.195627	1.0	1.0	0.0	0.25	0.25	1.00	1.00	...	0.0	0.0	1.00	1.00	0.50	0.50	0.50	0.50	-0.036118	0.0
393	1.0	0.0	0.059264	1.0	0.0	0.0	0.75	0.50	0.50	1.00	...	1.0	0.0	0.75	0.75	0.00	0.50	0.75	1.00	-0.076118	1.0
394	1.0	0.0	0.104718	0.0	0.0	0.0	0.25	0.25	1.00	0.75	...	1.0	0.0	0.50	0.25	0.50	0.50	0.50	1.00	-0.009451	0.0

395 rows × 31 columns

- Encoding:

```
def numerical_data():
    df['school'] = df['school'].map({'GP': 0, 'MS': 1})
    df['sex'] = df['sex'].map({'M': 0, 'F': 1})
    df['address'] = df['address'].map({'U': 0, 'R': 1})
    df['famsize'] = df['famsize'].map({'LE3': 0, 'GT3': 1})
    df['Pstatus'] = df['Pstatus'].map({'T': 0, 'A': 1})
    df['Mjob'] = df['Mjob'].map({'teacher': 0, 'health': 1, 'services': 2, 'at_home': 3, 'other': 4})
    df['Fjob'] = df['Fjob'].map({'teacher': 0, 'health': 1, 'services': 2, 'at_home': 3, 'other': 4})
    df['reason'] = df['reason'].map({'home': 0, 'reputation': 1, 'course': 2, 'other': 3})
    df['guardian'] = df['guardian'].map({'mother': 0, 'father': 1, 'other': 2})
    df['schoolsup'] = df['schoolsup'].map({'no': 0, 'yes': 1})
    df['famsup'] = df['famsup'].map({'no': 0, 'yes': 1})
    df['paid'] = df['paid'].map({'no': 0, 'yes': 1})
    df['activities'] = df['activities'].map({'no': 0, 'yes': 1})
    df['nursery'] = df['nursery'].map({'no': 0, 'yes': 1})
    df['higher'] = df['higher'].map({'no': 0, 'yes': 1})
    df['internet'] = df['internet'].map({'no': 0, 'yes': 1})
    df['romantic'] = df['romantic'].map({'no': 0, 'yes': 1})
    df['passed'] = df['passed'].map({'no': 0, 'yes': 1})
    # reorder dataframe columns :
    col = df['passed']
    del df['passed']
    df['passed'] = col
```

## SOURCE CODE & OUTPUT:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from time import time
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import confusion_matrix, roc_curve, accuracy_score, f1_score, roc_auc_score, classification_report
from astropy.table import Table
from sklearn.metrics import roc_auc_score

df = pd.read_csv('student.csv')
dfv = pd.read_csv('student.csv')
```

[ ] Python

df

[ ] Python Python

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	absences	passed
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	no	no	4	3	4	1	1	3	6	no
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	yes	no	5	3	3	1	1	3	4	no
2	GP	F	15	U	LE3	T	1	1	at_home	other	...	yes	no	4	3	2	2	3	3	10	yes
3	GP	F	15	U	GT3	T	4	2	health	services	...	yes	yes	3	2	2	1	1	5	2	yes
4	GP	F	16	U	GT3	T	3	3	other	other	...	no	no	4	3	2	1	2	5	4	yes
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
390	MS	M	20	U	LE3	A	2	2	services	services	...	no	no	5	5	4	4	5	4	11	no
391	MS	M	17	U	LE3	T	3	1	services	services	...	yes	no	2	4	5	3	4	2	3	yes

```
def numerical_data():
    df['school'] = df['school'].map({'GP': 0, 'MS': 1})
    df['sex'] = df['sex'].map({'M': 0, 'F': 1})
    df['address'] = df['address'].map({'U': 0, 'R': 1})
    df['famsize'] = df['famsize'].map({'LE3': 0, 'GT3': 1})
    df['Pstatus'] = df['Pstatus'].map({'T': 0, 'A': 1})
    df['Mjob'] = df['Mjob'].map({'teacher': 0, 'health': 1, 'services': 2, 'at_home': 3, 'other': 4})
    df['Fjob'] = df['Fjob'].map({'teacher': 0, 'health': 1, 'services': 2, 'at_home': 3, 'other': 4})
    df['reason'] = df['reason'].map({'home': 0, 'reputation': 1, 'course': 2, 'other': 3})
    df['guardian'] = df['guardian'].map({'mother': 0, 'father': 1, 'other': 2})
    df['schoolsup'] = df['schoolsup'].map({'no': 0, 'yes': 1})
    df['famsup'] = df['famsup'].map({'no': 0, 'yes': 1})
    df['paid'] = df['paid'].map({'no': 0, 'yes': 1})
    df['activities'] = df['activities'].map({'no': 0, 'yes': 1})
    df['nursery'] = df['nursery'].map({'no': 0, 'yes': 1})
    df['higher'] = df['higher'].map({'no': 0, 'yes': 1})
    df['internet'] = df['internet'].map({'no': 0, 'yes': 1})
    df['romantic'] = df['romantic'].map({'no': 0, 'yes': 1})
    df['passed'] = df['passed'].map({'no': 0, 'yes': 1})
    # reorder dataframe columns :
    col = df['passed']
    del df['passed']
    df['passed'] = col

    # feature scaling will allow the algorithm to converge faster, large data will have same scal

def feature_scaling(df):
    for i in df:
        col = df[i]
        # let's choose columns that have large values
        if(np.max(col)>6):
            Max = max(col)
            Min = min(col)
            mean = np.mean(col)
            col = (col-mean)/(Max)
            df[i] = col
```

```
# All values in numerical after calling numerical_data() function
numerical_data()
df
```

Python

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	absences	passed
0	0	1	18	0	1	1	4	4	3	0	...	0	0	4	3	4	1	1	3	6	0
1	0	1	17	0	1	0	1	1	3	4	...	1	0	5	3	3	1	1	3	4	0
2	0	1	15	0	0	0	1	1	3	4	...	1	0	4	3	2	2	3	3	10	1
3	0	1	15	0	1	0	4	2	1	2	...	1	1	3	2	2	1	1	5	2	1
4	0	1	16	0	1	0	3	3	4	4	...	0	0	4	3	2	1	2	5	4	1
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
390	1	0	20	0	0	1	2	2	2	2	...	0	0	5	5	4	4	5	4	11	0
391	1	0	17	0	0	0	3	1	2	2	...	1	0	2	4	5	3	4	2	3	1
392	1	0	21	1	1	0	1	1	4	4	...	0	0	5	5	3	3	3	3	3	0
393	1	0	18	1	0	0	3	2	2	4	...	1	0	4	4	1	3	4	5	0	1
394	1	0	19	0	0	0	1	1	4	3	...	1	0	3	2	3	3	3	5	5	0

395 rows × 31 columns

```
# Let's scal our features
feature_scaling(df)

# Now we are ready for models training
df
```

Python

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	absences	passed
0	0.0	1.0	0.059264	0.0	1.0	1.0	1.00	1.00	0.75	0.00	...	0.0	0.0	0.75	0.50	0.75	0.00	0.00	0.50	0.003882	0.0
1	0.0	1.0	0.013809	0.0	1.0	0.0	0.25	0.25	0.75	1.00	...	1.0	0.0	1.00	0.50	0.50	0.00	0.00	0.50	-0.022785	0.0
2	0.0	1.0	-0.077100	0.0	0.0	0.0	0.25	0.25	0.75	1.00	...	1.0	0.0	0.75	0.50	0.25	0.25	0.50	0.50	0.057215	1.0
3	0.0	1.0	-0.077100	0.0	1.0	0.0	1.00	0.50	0.25	0.50	...	1.0	1.0	0.50	0.25	0.25	0.00	0.00	1.00	-0.049451	1.0
4	0.0	1.0	-0.031646	0.0	1.0	0.0	0.75	0.75	1.00	1.00	...	0.0	0.0	0.75	0.50	0.25	0.00	0.25	1.00	-0.022785	1.0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
390	1.0	0.0	0.150173	0.0	0.0	1.0	0.50	0.50	0.50	0.50	...	0.0	0.0	1.00	1.00	0.75	0.75	1.00	0.75	0.070549	0.0
391	1.0	0.0	0.013809	0.0	0.0	0.0	0.75	0.25	0.50	0.50	...	1.0	0.0	0.25	0.75	1.00	0.50	0.75	0.25	-0.036118	1.0
392	1.0	0.0	0.195627	1.0	1.0	0.0	0.25	0.25	1.00	1.00	...	0.0	0.0	1.00	1.00	0.50	0.50	0.50	0.50	-0.036118	0.0
393	1.0	0.0	0.059264	1.0	0.0	0.0	0.75	0.50	0.50	1.00	...	1.0	0.0	0.75	0.75	0.00	0.50	0.75	1.00	-0.076118	1.0
394	1.0	0.0	0.104718	0.0	0.0	0.0	0.25	0.25	1.00	0.75	...	1.0	0.0	0.50	0.25	0.50	0.50	0.50	1.00	-0.009451	0.0

395 rows × 31 columns

df.shape

Python

(395, 31)

```

df.dropna().shape # their is no null value "fortunately:"
[ ]

... (395, 31)

df.columns
[ ]

Index(['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
      'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
      'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
      'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
      'Walc', 'health', 'absences', 'passed'],
      dtype=object)

features=['school', 'sex', 'age', 'address', 'famsize', 'Pstatus', 'Medu', 'Fedu',
          'Mjob', 'Fjob', 'reason', 'guardian', 'traveltime', 'studytime',
          'failures', 'schoolsup', 'famsup', 'paid', 'activities', 'nursery',
          'higher', 'internet', 'romantic', 'famrel', 'freetime', 'goout', 'Dalc',
          'Walc', 'health', 'absences']
[ ]

#plot of student status
dfv['passed'].value_counts()
[ ]


... yes    265
    no     130
Name: passed, dtype: int64

```

```

labels = 'student pass the final exam ', 'student fail the final exam'
sizes = [265, 130]
colors=['lightskyblue','yellow']
fig1, ax1 = plt.subplots()
ax1.pie(sizes, labels=labels, autopct='%1.1f%%',colors=colors,
        shadow=True, startangle=90)
ax1.axis('equal') # Equal aspect ratio ensures that pie is drawn as a circle.
plt.show()


```



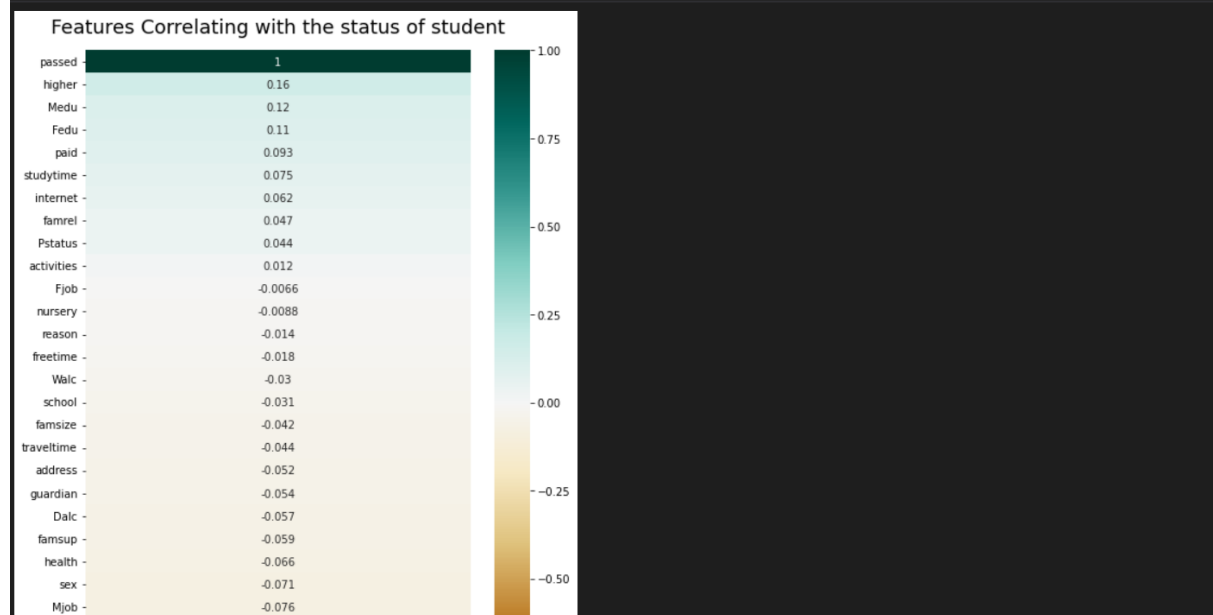
```
# see correlation between variables through a correlation heatmap
corr = df.corr()
plt.figure(figsize=(30,30))
sns.heatmap(corr, annot=True, cmap="Reds")
plt.title('Correlation Heatmap', fontsize=20)
```

Python

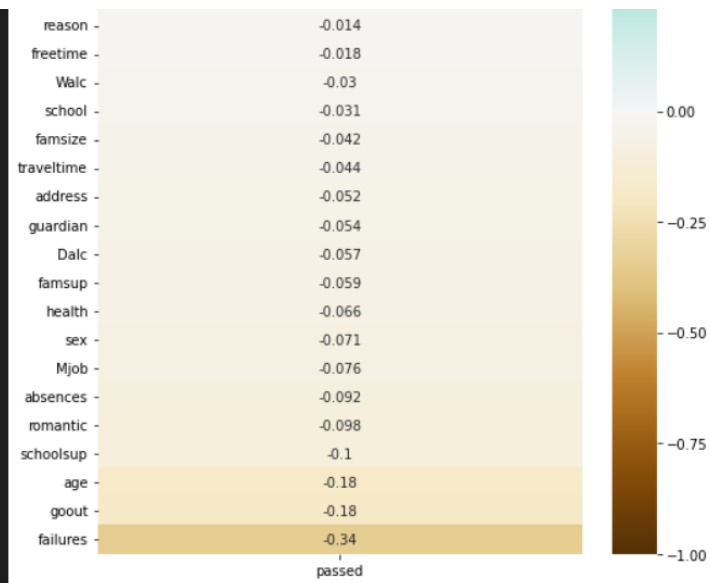
Text(0.5, 1.0, 'Correlation Heatmap')



```
plt.figure(figsize=(8, 12))
heatmap = sns.heatmap(df.corr()[['passed']].sort_values(by='passed', ascending=False), vmin=-1, vmax=1, annot=True, cmap='BrBG')
heatmap.set_title('Features Correlating with the status of student', fontdict={'fontsize':18}, pad=16);
```








```
df["goout"].unique()
```

```
array([0.75, 0.5 , 0.25, 0. , 1.  ])
```

```
# going out
perc = (lambda col: col/col.sum())
index = [0,1]
out_tab = pd.crosstab(index=df.passed, columns=df.goout)
out_perc = out_tab.apply(perc).reindex(index)
out_perc.plot.bar(colormap="mako_r", fontsize=16, figsize=(14,6))
plt.title('student status By Frequency of Going Out', fontsize=20)
plt.ylabel('Percentage of Student', fontsize=16)
plt.xlabel('Student status', fontsize=16)
```

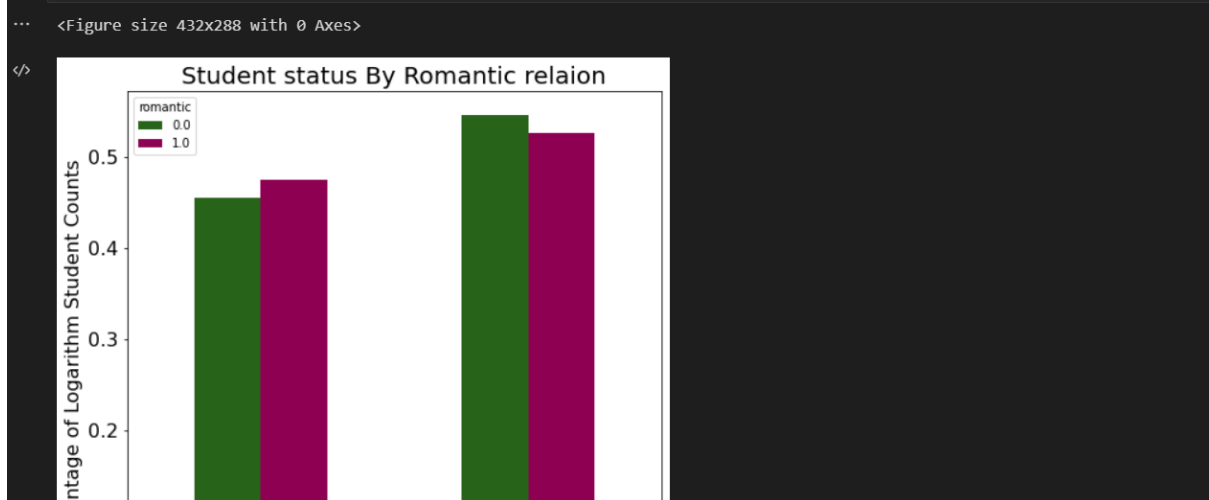
```
Text(0.5, 0, 'Student status')
```



```

# romantic status
romance_tab1 = pd.crosstab(index=df.passed, columns=df.romantic)
romance_tab = np.log(romance_tab1)
romance_perc = romance_tab.apply(perc).reindex(index)
plt.figure()
romance_perc.plot.bar(colormap="PiYG_r", fontsize=16, figsize=(8,8))
plt.title('Student status By Romantic relaion', fontsize=20)
plt.ylabel('Percentage of Logarithm Student Counts ', fontsize=16)
plt.xlabel('Student status', fontsize=16)
plt.show()
# 0 in romantic mean no romantic relation


```



```

# 1) mother job
# Mjob distribution
f, fx = plt.subplots()
figure = sns.countplot(x = 'Mjob', data=dfv, order=['teacher', 'health', 'services', 'at_home', 'other'])
fx = fx.set(ylabel="Count", xlabel="Mother Job")
figure.grid(False)

```




```
mjob_perc.plot.bar(colormap="mako_r", fontsize=16, figsize=(8,8))
plt.title('Student status By mother JOB', fontsize=20)
plt.ylabel('Percentage of Logarithm Student Counts ', fontsize=16)
plt.xlabel('Student status', fontsize=16)
plt.show()
# 'teacher': 0, 'health': 1, 'services': 2, 'at_home': 3, 'other': 4
```

[ ]

... <Figure size 432x288 with 0 Axes>

</>




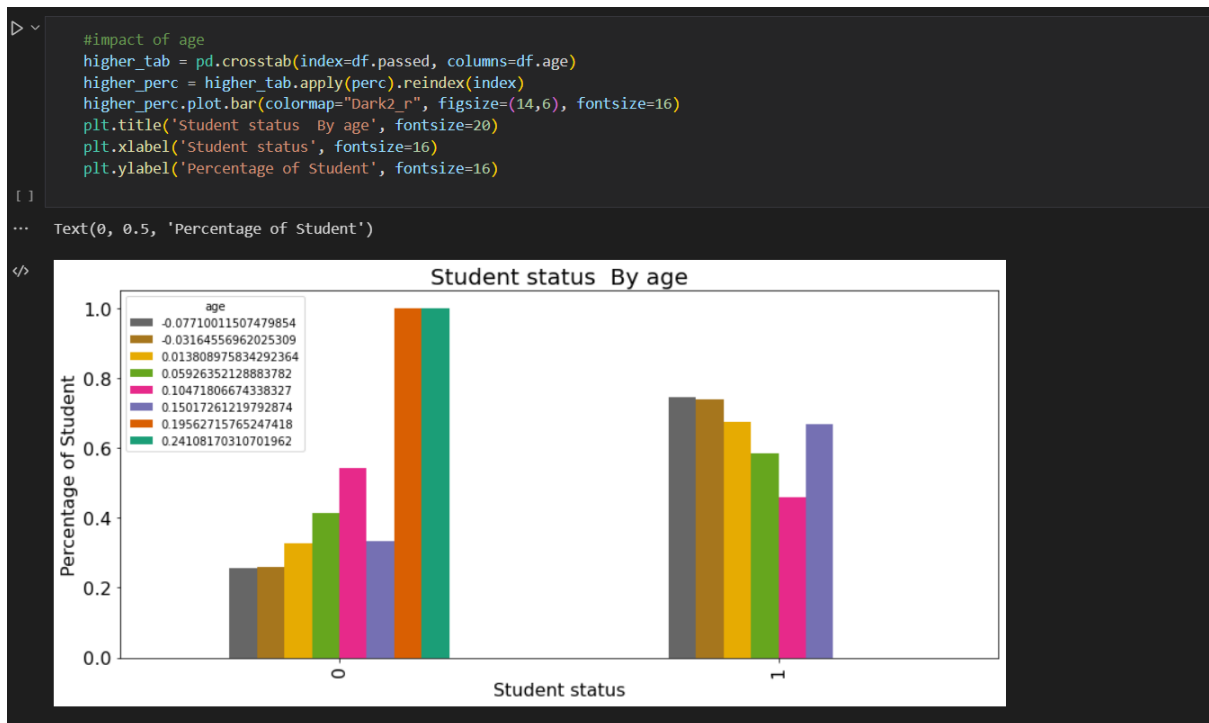
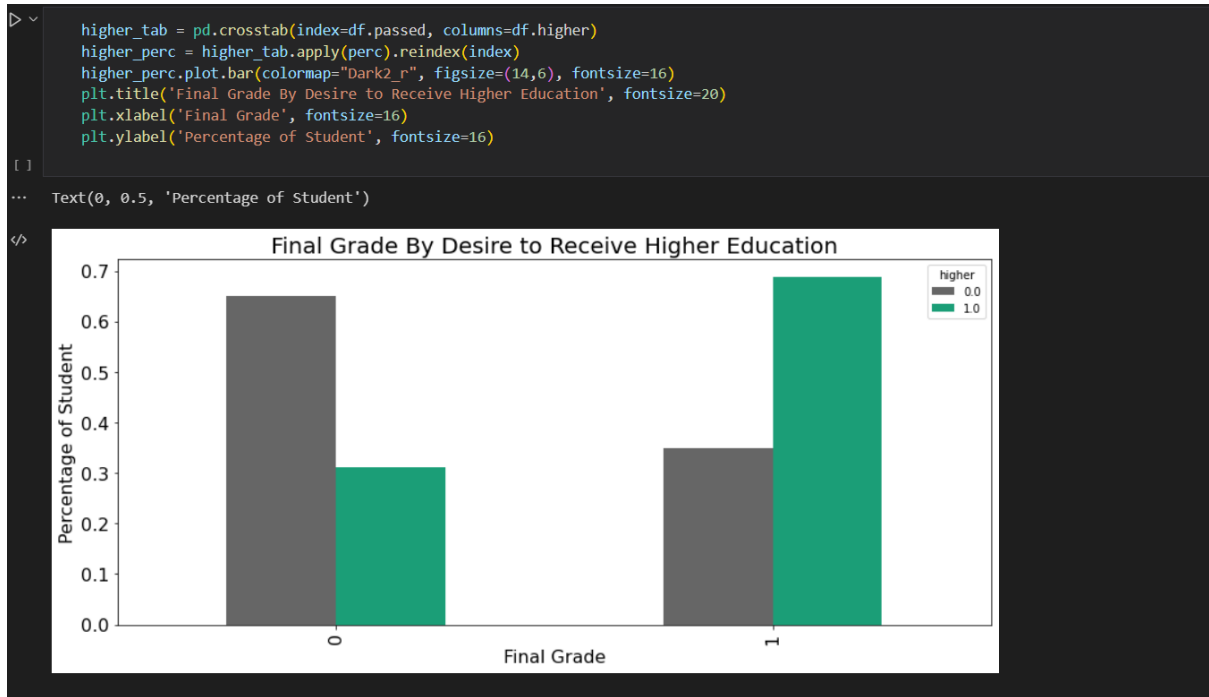
▷ ▾

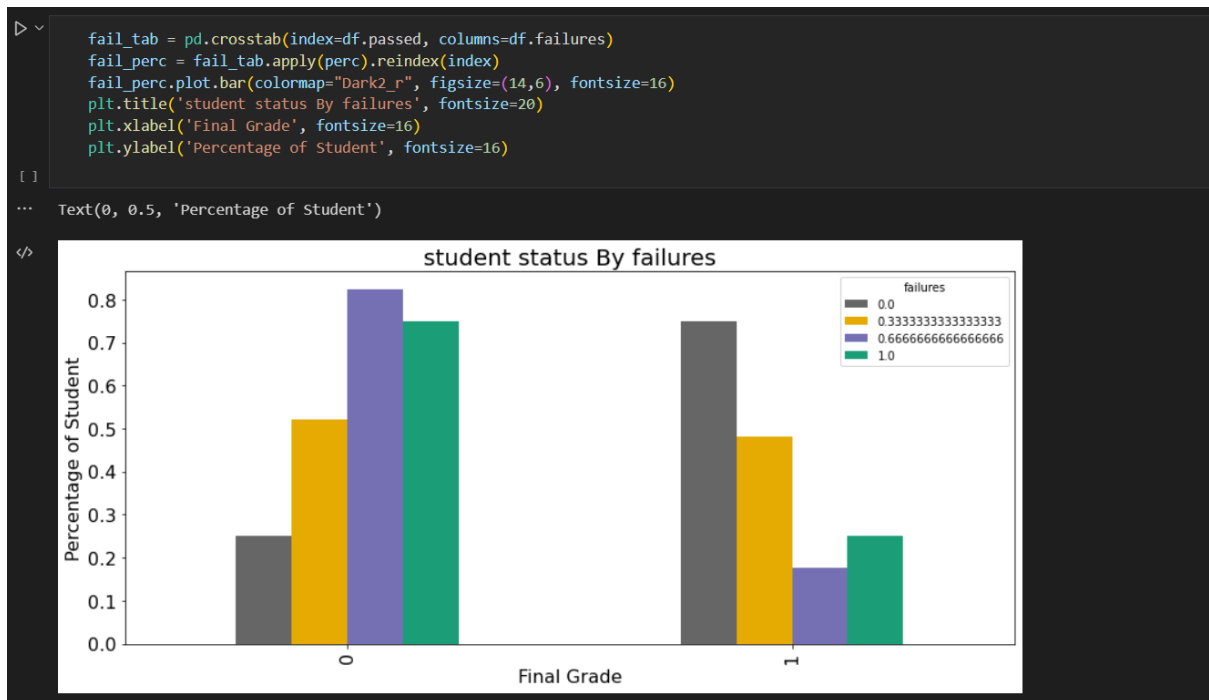
```
#Mother education:
good = df.loc[df.passed==1]
poor=df.loc[df.passed==0]
good['good_student_mother_education'] = good.Medu
poor['poor_student_mother_education'] = poor.Medu
plt.figure(figsize=(6,4))
p=sns.kdeplot(good['good_student_mother_education'], shade=True, color="r")#good_student in red
p=sns.kdeplot(poor['poor_student_mother_education'], shade=True, color="b")#poor_student in blue
plt.xlabel('Mother Education Level', fontsize=20)
```

[ ]

</>







```


ad_tab1 = pd.crosstab(index=df.passed, columns=df.address)
ad_tab = np.log(ad_tab1)
ad_perc = ad_tab.apply(perc).reindex(index)
ad_perc.plot.bar(colormap="RdYlGn_r", fontsize=16, figsize=(8,6))
plt.title('student status By Living Area', fontsize=20)
plt.ylabel('Percentage of Logarithm Student#', fontsize=16)
plt.xlabel('Student status', fontsize=16)

```

[ ]

... Text(0.5, 0, 'Student status')

</>



```


#impact of weekend alcohol consumption in student performance
alc_tab = pd.crosstab(index=df.passed, columns=df.Walc)
alc_perc = alc_tab.apply(perc).reindex(index)
alc_perc.plot.bar(colormap="Dark2_r", figsize=(14,6), fontsize=16)
plt.title('student status By weekend alchol consumption', fontsize=20)
plt.xlabel('Student status', fontsize=16)
plt.ylabel('Percentage of Student', fontsize=16)

```


[ ]

... Text(0, 0.5, 'Percentage of Student')

</>

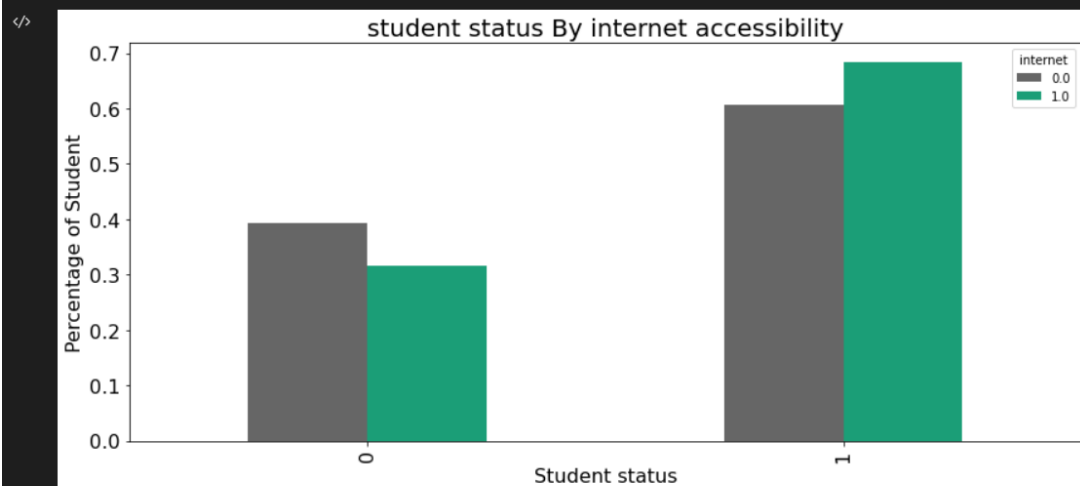


```
# weekend alcohol consumption
# create good student dataframe
good = df.loc[df.passed == 1]
good['good_alcohol_usage']=good.walc
# create poor student dataframe
poor = df.loc[df.passed == 0]
poor['poor_alcohol_usage']=poor.walc
plt.figure(figsize=(10,6))
p1=sns.kdeplot(good['good_alcohol_usage'], shade=True, color="r")
p1=sns.kdeplot(poor['poor_alcohol_usage'], shade=True, color="b")
plt.title('Good Performance vs. Poor Performance Student Weekend Alcohol Consumption', fontsize=20)
plt.ylabel('Density', fontsize=16)
plt.xlabel('Level of Alcohol Consumption', fontsize=16)
```



```
alc_tab = pd.crosstab(index=df.passed, columns=df.internet)
alc_perc = alc_tab.apply(perc).reindex(index)
alc_perc.plot.bar(colormap="Dark2_r", figsize=(14,6), fontsize=16)
plt.title('student status By internet accessibility', fontsize=20)
plt.xlabel('Student status', fontsize=16)
plt.ylabel('Percentage of Student', fontsize=16)
```

```
[ ]
... Text(0, 0.5, 'Percentage of Student')
```



```


stu_tab = pd.crosstab(index=df.passed, columns=df.studytime)
stu_perc = stu_tab.apply(perc).reindex(index)
stu_perc.plot.bar(colormap="Dark2_r", figsize=(14,6), fontsize=16)
plt.title('student status By study time', fontsize=20)
plt.xlabel('Student status', fontsize=16)
plt.ylabel('Percentage of Student', fontsize=16)

```

[ ]

... Text(0, 0.5, 'Percentage of Student')

</>



```


he_tab = pd.crosstab(index=df.passed, columns=df.health)
he_perc = he_tab.apply(perc).reindex(index)
he_perc.plot.bar(colormap="Dark2_r", figsize=(14,6), fontsize=16)
plt.title('student status By health', fontsize=20)
plt.xlabel('Student status', fontsize=16)
plt.ylabel('Percentage of Student', fontsize=16)

```

[ ]

... Text(0, 0.5, 'Percentage of Student')

</>





```

# split data train 70 % and test 30 %

data = df.to_numpy()
n = data.shape[1]
x = data[:,0:n-1]
y = data[:,n-1]
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=0)

# Once our data is split, we can forget about x_test and y_test until we define our model.
#x_train and y_train are the samples we will use to train the model

1

# let's create a model and train it

logisticRegr = LogisticRegression(C=1)

]

#and now let's do the training

logisticRegr.fit(x_train,y_train)

1

LogisticRegression(C=1)

```

```

y_pred=logisticRegr.predict(x_test)
y_pred

[ ]

... array([1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0.,
        1., 1., 1., 1., 0., 1., 1., 1., 1., 1., 1., 1., 1., 0., 0., 1., 0.,
        1., 0., 1., 1., 1., 1., 1., 1., 0., 0., 1., 0., 1., 1., 1., 0., 1.,
        1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0.,
        1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 0., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1., 1.,
        1., 0., 1., 1., 1., 1., 0., 0., 1., 1., 1., 1., 0., 1., 1., 1.,
        1., 0., 1., 1., 1., 1., 0., 0., 1., 1., 1., 1., 0., 1., 1., 1., 1.])

```

```

#let's have a look at the accuracy of the model

Sctest=logisticRegr.score(x_test,y_test)
Sctrain=logisticRegr.score(x_train,y_train)

print('#Accuracy test is: ',Sctest)
print('#Accuracy train is: ',Sctrain)

f1 = f1_score(y_test, y_pred, average='macro')

print('\n#f1 score is: ',f1)

[ ]

... #Accuracy test is:  0.6386554621848739
#Accuracy train is:  0.7463768115942029

#f1 score is:  0.5533734834598935

```

```
#let's have a look at the accuracy of the model

Sctest=logisticRegr.score(x_test,y_test)
Sctrain=logisticRegr.score(x_train,y_train)

print('Accuracy test is: ',Sctest)
print('Accuracy train is: ',Sctrain)
```

[ ]

```
... Accuracy test is:  0.6386554621848739
Accuracy train is:  0.7463768115942029
```

▷ ▾

```
#now, we can get the confusion matrix with confusion_matrix():

confusion_matrix(y_test, y_pred)
```

[ ]

```
... array([[12, 38],
          [ 5, 64]])
```

+ Code + Markdown



```
#import classification_report
print(classification_report(y_test, y_pred))
```

[ ]

...


	precision	recall	f1-score	support
0.0	0.71	0.24	0.36	50
1.0	0.63	0.93	0.75	69
accuracy			0.64	119
macro avg	0.67	0.58	0.55	119
weighted avg	0.66	0.64	0.58	119

```
#ploting the roc_curve
```

```
fpositif, tpositif, thresholds = roc_curve(y_test, y_pred)
plt.plot([0,1],[0,1], '--')
plt.plot(fpositif, tpositif, label='LogisticRegr')
plt.xlabel('false positif')
plt.ylabel('true positif')
plt.title('LogisticRegr ROC curve')
p=plt.show()
```

[ ]

...



```

max_iteration = 0
maxF1 = 0
maxAccuracy = 0
optimal_state = 0
import random
for k in range(max_iteration):
    print ('Iteration :'+str(k)+' , Current accuracy: '+str(maxAccuracy)+ ' , Current f1 : '+str(maxF1), end="\r")
    split_state = np.random.randint(1,100000000)-1
    x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=split_state)
    logisticRegr = LogisticRegression(C=1)
    logisticRegr.fit(x_train,y_train)
    y_pred=logisticRegr.predict(x_test)
    f1 = f1_score(y_test, y_pred, average='macro')
    accuracy = accuracy_score(y_test, y_pred)*100

    if (accuracy>maxAccuracy and f1>maxF1):
        maxF1 = f1
        maxAccuracy = accuracy
        optimal_state = split_state

optimal_state = 85491961
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=optimal_state)
logisticRegr = LogisticRegression(C=1)
logisticRegr.fit(x_train,y_train)
y_pred=logisticRegr.predict(x_test)
f1 = f1_score(y_test, y_pred, average='macro')
accuracy = accuracy_score(y_test, y_pred)*100
print('\n\n*Accuracy is: '+str(accuracy)+'\n*f1 score is: ',f1)

yt_lg,yp_lg = y_test,y_pred
#ploting the roc_curve

print ( '\n\n *the ROC curve: ')

```

8 Ju

```

fpositif, tpositif, thresholds = roc_curve(y_test, y_pred)
plt.plot([0,1],[0,1], '--')
plt.plot(fpositif, tpositif, label='LogisticRegr')
plt.xlabel('false positif')
plt.ylabel('true positif')
plt.title('LogisticRegr ROC curve')
p=plt.show()

```

#visualizing the confusion matrix:

```
print ( ' *the confusion matrix ' )
```

```
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm,annot=True)
```

[ ]


...

\*Accuracy is: 80.67226890756302

\*f1 score is: 0.7408389357068459

\*the ROC curve:

</>



\*the confusion matrix

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fa8997bf050>

</>



```
#define data
y=df.passed
target=["passed"]
x = df.drop(target,axis = 1 )
```

```
max_iteration = 0
maxF1 = 0
maxAccuracy = 0
optimal_state = 0
for k in range(max_iteration):
    print ('Iteration :'+str(k)+', Current accuracy: '+str(maxAccuracy)+ ', Current f1 : '+str(maxF1), end="\n")
    split_state = np.random.randint(1,100000000)-1
    x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=split_state)
    KNN = KNeighborsClassifier()
    KNN.fit(x_train,y_train)
    y_pred=KNN.predict(x_test)
    f1 = f1_score(y_test, y_pred, average='macro')
    accuracy = accuracy_score(y_test, y_pred)*100

    if (accuracy>maxAccuracy and f1>maxF1):
        maxF1 = f1
        maxAccuracy = accuracy
        optimal_state = split_state

optimal_state = 71027464
```

```

x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=optimal_state)
KNN= KNeighborsClassifier()
KNN.fit(x_train,y_train)
y_pred=KNN.predict(x_test)
f1 = f1_score(y_test, y_pred, average='macro')
accuracy = accuracy_score(y_test, y_pred)*100
print('\n\n*Accuracy is: '+str(accuracy)+'\n*f1 score is: ',f1)

print ('random_state is ',optimal_state)

#ploting the roc_curve

print ( '\n\n *the ROC curve: ')

fpositif, tpositif, thresholds = roc_curve(y_test, y_pred)
plt.plot([0,1],[0,1], '--')
plt.plot(fpositif, tpositif, label='knn')
plt.xlabel('false positif')
plt.ylabel('true positif')
plt.title('KNN ROC curve')
p=plt.show()

yt_knn,yp_knn= y_test,y_pred
#visualizing the confusion matrix:

print (' *the confusion matrix ')

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm,annot=True)

```

```

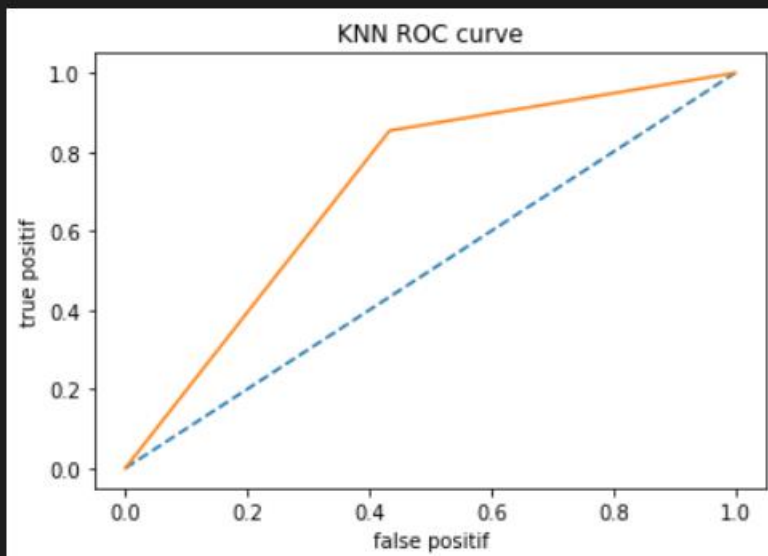
*Accuracy is: 78.15126050420169
*f1 score is:  0.7102996254681648
random_state is  71027464

```

```


*the ROC curve:

```



\*the confusion matrix

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fa899538f10>



```
#Setup arrays to store training and test accuracies
neighbors= np.arange(1,20)
train_accuracy =np.empty(19)
test_accuracy = np.empty(19)


for i,k in enumerate(neighbors):
    #Setup a knn classifier with k neighbors
    knn = KNeighborsClassifier(n_neighbors=k)

    #Fit the model
    knn.fit(x_train, y_train)

    #Compute accuracy on the training set
    train_accuracy[i] = knn.score(x_train, y_train)

    #Compute accuracy on the test set
    test_accuracy[i] = knn.score(x_test, y_test)

# Plotting the curve
plt.title('k-NN Varying number of neighbors')
plt.plot(neighbors, test_accuracy, label='Testing Accuracy')
plt.plot(neighbors, train_accuracy, label='Training accuracy')
plt.legend()
plt.xlabel('Number of neighbors')
plt.ylabel('Accuracy')
plt.show()
```



```
#In case of classifier like knn the parameter to be tuned is n_neighbors
param_grid = {'n_neighbors':np.arange(1,20)}
knn = KNeighborsClassifier()
knn_cv= GridSearchCV(knn,param_grid,cv=5)
knn_cv.fit(x_train,y_train)
#best score\n",
knn_cv.best_score_
```

[ ]

0.6449350649350649

knn\_cv.best\_params\_

[ ]

{'n\_neighbors': 19}

```
param_grid = {'n_neighbors':np.arange(1,20)}
knn = KNeighborsClassifier()
knn_cv= GridSearchCV(knn,param_grid,cv=5)
knn_cv.fit(x,y)
#best score\n",
knn_cv.best_score_
```

[ ]

0.6734177215189873

knn\_cv.best\_params\_

[ ]

{'n\_neighbors': 7}



```

params = {"n_neighbors": [7, 19], "metric": ["euclidean", "manhattan", "chebyshev"]}
acc = {}

```

```

for m in params["metric"]:
    acc[m] = []
    for k in params["n_neighbors"]:
        print("Model_{} metric: {}, n_neighbors: {}".format(i, m, k))
        i += 1
        t = time()
        knn = KNeighborsClassifier(n_neighbors=k, metric=m)
        knn.fit(x_train, y_train)
        pred = knn.predict(x_test)
        print("Time: ", time() - t)
        acc[m].append(accuracy_score(y_test, y_pred))
        print("Acc: ", acc[m][-1])

```

[ ]

```

... Model_18 metric: euclidean, n_neighbors: 7
Time: 0.012510061264038086
Acc: 0.7815126050420168
Model_19 metric: euclidean, n_neighbors: 19
Time: 0.011599302291870117
Acc: 0.7815126050420168
Model_20 metric: manhattan, n_neighbors: 7
Time: 0.012638568878173828
Acc: 0.7815126050420168
Model_21 metric: manhattan, n_neighbors: 19
Time: 0.012667417526245117
Acc: 0.7815126050420168
Model_22 metric: chebyshev, n_neighbors: 7
Time: 0.011996269226074219
Acc: 0.7815126050420168

```

▷ ▾

```

max_iteration = 0
maxF1 = 0
maxAccuracy = 0
optimal_state = 0
f1 = 0
accuracy = 0
True60 = False
for k in range(max_iteration):
    print('Iteration :'+str(k)+' , Current accuracy: '+str(maxAccuracy)+ ' , Current f1 : '+str(maxF1), end="\n")
    split_state = np.random.randint(1,100000000)-1
    x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=split_state)
    KNN = KNeighborsClassifier(n_neighbors=7,metric='chebyshev')
    KNN.fit(x_train,y_train)
    y_pred=KNN.predict(x_test)
    f1 = f1_score(y_test, y_pred, average='macro')
    accuracy = accuracy_score(y_test, y_pred)*100

    if accuracy>maxAccuracy and f1>=0.5:
        maxF1 = f1
        maxAccuracy = accuracy
        optimal_state = split_state
        if maxAccuracy>79:
            break

optimal_state = 29300362
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_state=optimal_state)
KNN_f= KNeighborsClassifier(n_neighbors=7,metric='chebyshev')
KNN_f.fit(x_train,y_train)
y_pred=KNN_f.predict(x_test)
f1 = f1_score(y_test, y_pred, average='macro')
accuracy = accuracy_score(y_test, y_pred)*100
print('\n\n*Accuracy is: '+str(accuracy)+'\n*f1 score is: ',f1)

```


```
yt_knn,yp_knn= y_test,y_pred

[ ]
...

*Accuracy is: 69.74789915966386
*f1 score is: 0.47959183673469385
random_state is 29300362
```

```
▷ ▾
ac = accuracy_score(yt_knn,yp_knn)
print('Accuracy is: ',ac)
cm= confusion_matrix(yt_knn,yp_knn)
sns.heatmap(cm,annot=True)
yt_knn,yp_knn = y_test,y_pred
```

```
[ ]
... Accuracy is: 0.6974789915966386
```



```
print(classification_report(y_test,y_pred))

[ ]
...

```

	precision	recall	f1-score	support
0.0	0.19	0.12	0.14	26
1.0	0.78	0.86	0.82	93
accuracy			0.70	119
macro avg	0.48	0.49	0.48	119
weighted avg	0.65	0.70	0.67	119

```
▷ ▾
#ploting the roc_curve


print ( ' the ROC curve: ' )

fpositif, tpositif, thresholds = roc_curve(y_test, y_pred)
plt.plot([0,1],[0,1], '--')
plt.plot(fpositif, tpositif, label='final knn model')
plt.xlabel('false positif')
plt.ylabel('true positif')
plt.title('knn_f ROC curve')
p=plt.show()
```

```
[ ]
```

```
... the ROC curve:
```

```
/>
```



```
def showResults(accuracy, trainingTime, y_pred,model):

    print('-----Results :',model,'-----')
    confusionMatrix = confusion_matrix(y_test, y_pred)
    print('\n The ROC curve is :\n')
    fig, _ = plt.subplots()
    fpr,tpr,thresholds=roc_curve(y_test,y_pred)
    plt.plot([0, 1],[0, 1], '--')
    plt.plot(fpr,tpr,label=model)
    plt.xlabel('false positive')
    plt.ylabel('false negative')
    plt.legend()
    fig.suptitle('ROC curve: '+str(model))
    plt.show()

    print('-----')
    print('The model accuracy:', round(accuracy), '%')
    print('-----')
    print('The training time is: ',trainingTime)
    print('-----')
    print('The f1 score is :',round(100*f1_score(y_test, y_pred, average='macro'))/100)
    print('-----')
    print('The roc_auc_score is :',round(100*roc_auc_score(y_test, y_pred))/100)
    print('-----')
    print('The confusion matrix is :\n')
    ax = plt.axes()
    sns.heatmap(confusionMatrix,annot=True)
```

```

# Optimal C
def optimal_C_value():
    Ci = np.array(( 0.0001,0.001,0.01,0.05,0.1,4,10,40,100))
    minError = float('Inf')
    optimal_C = float('Inf')

    for c in Ci:
        clf = SVC(C=c,kernel='linear')
        clf.fit(X_train, y_train)
        predictions = clf.predict(X_val)
        error = np.mean(np.double(predictions != y_val))
        if error < minError:
            minError = error
            optimal_C = c
    return optimal_C

# Optimal C and the degree of the polynomial
def optimal_C_d_values():
    Ci = np.array(( 0.0001,0.001,0.01,0.05,0.1,4,10,40,100))
    Di = np.array(( 2, 5, 10, 15, 20, 25, 30))
    minError = float('Inf')
    optimal_C = float('Inf')
    optimal_d = float('Inf')

    for d in Di:
        for c in Ci:
            clf = SVC(C=c,kernel='poly', degree=d)
            clf.fit(X_train, y_train)
            predictions = clf.predict(X_val)
            error = np.mean(np.double(predictions != y_val))
            if error < minError:
                minError = error
                optimal_C = c
                optimal_d = d
    return optimal_C,optimal_d

```

```

# Optimal C and gamma
def optimal_C_gamma_values():
    Ci = np.array(( 0.0001,0.001,0.01,0.05,0.1,4,10,40,100))
    Gi = np.array(( 0.000001,0.00001,0.01,1,2,3,5,20,70,100,500,1000))
    minError = float('Inf')
    optimal_C = float('Inf')
    optimal_g = float('Inf')

    for g in Gi:
        for c in Ci:
            clf = SVC(C=c,kernel='rbf', gamma=g)
            clf.fit(X_train, y_train)
            predictions = clf.predict(X_val)
            error = np.mean(np.double(predictions != y_val))
            if error < minError:
                minError = error
                optimal_C = c
                optimal_g = g
    return optimal_C,optimal_g

# -----
# Compare the three kernels

def compare_kernels():
    X_train1,X_val1,X_test1,y_train1,y_val1,y_test1 = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state1)
    X_train2,X_val2,X_test2,y_train2,y_val2,y_test2 = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state2)
    X_train3,X_val3,X_test3,y_train3,y_val3,y_test3 = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state3)
    print('----- Comparison -----')
    print('\n')
    f11 = "{:.2f}".format(f1_score(y_test1, y_linear, average='macro'))
    f22 = "{:.2f}".format(f1_score(y_test2, y_poly, average='macro'))
    f33 = "{:.2f}".format(f1_score(y_test3, y_gauss, average='macro'))
    roc1 = "{:.2f}".format(roc_auc_score(y_test1, y_linear))
    roc2 = "{:.2f}".format(roc_auc_score(y_test2, y_poly))
    roc3 = "{:.2f}".format(roc_auc_score(y_test3, y_gauss))

```

```

a1,a2 = confusion_matrix(y_test1, y_linear)[0],confusion_matrix(y_test1, y_linear)[1]
b1,b2 = confusion_matrix(y_test2, y_poly)[0],confusion_matrix(y_test2, y_poly)[1]
c1,c2 = confusion_matrix(y_test3, y_gauss)[0],confusion_matrix(y_test3, y_gauss)[1]
data_rows = [('training time',time1, time2, time3),
              ('', '', '', ''),
              ('accuracy %',linear_accuracy, poly_accuracy, gauss_accuracy),
              ('', '', '', ''),
              ('confusion matrix',a1, b1, c1),
              ('',a2,b2,c2),
              ('', '', '', ''),
              ('f1 score',f11,f22,f33),
              ('', '', '', ''),
              ('roc_auc_score',roc1,roc2,roc3)]
t = Table(rows=data_rows, names=('metric', 'Linear kernel', 'polynomial kernel', 'gaussian kernel'))
print(t)
print('\n\n')
print('The Roc curves :\n')
y_pred1 = y_linear
y_pred2 = y_poly
y_pred3 = y_gauss
fig, _ = plt.subplots()
fig.suptitle('Comparison of three ROC curves')
fpr,tpr,thresholds=roc_curve(y_test1,y_pred1)
plt.plot([0, 1],[0, 1], '--')
plt.plot(fpr,tpr,label='Linear kernel :'+str(roc1))
plt.xlabel('false positive')
plt.ylabel('false negative')
fpr,tpr,thresholds=roc_curve(y_test2,y_pred2)
plt.plot(fpr,tpr,label='Polynomial kernel :'+str(roc2))
fpr,tpr,thresholds=roc_curve(y_test3,y_pred3)
plt.plot(fpr,tpr,label='Gaussian kernel :'+str(roc3))
plt.legend()
plt.show()

```

```

def best_kernel(kernel):
    x_train1,X_val1,X_test1,y_train1,y_val1,y_test1 = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state1)
    x_train2,X_val2,X_test2,y_train2,y_val2,y_test2 = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state2)
    x_train3,X_val3,X_test3,y_train3,y_val3,y_test3 = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state3)

    time = 0
    f1 = 0
    accuracy = 0
    rc = 0
    y = 0
    if kernel == 'linear kernel':
        time = time1
        f1 = "{:.2f}".format(f1_score(y_test1, y_linear, average='macro'))
        accuracy = round(100*linear_accuracy)/100
        rc = round(100*roc_auc_score(y_test1, y_linear))/100
        y_test = y_test1
        y = y_linear
    elif kernel == 'polynomial kernel':
        time = time2
        f1 = "{:.2f}".format(f1_score(y_test2, y_poly, average='macro'))
        accuracy = round(100*poly_accuracy)/100
        rc = round(100*roc_auc_score(y_test2, y_poly))/100
        y_test = y_test2
        y = y_poly
    else :
        time = time3
        f1 = "{:.2f}".format(f1_score(y_test3, y_gauss, average='macro'))
        accuracy = round(100*gauss_accuracy)/100
        rc = round(100*roc_auc_score(y_test3, y_gauss))/100
        y_test = y_test3
        y = y_gauss

    # used for comparing three classifiers(knn, logistic regression and svm)
    yt_svm,yp_svm = y_test, y

```

```

print('The choosen kernel :',kernel)
print('the training :',time)
print('the accuracy :',round(accuracy),'%')
print('the f1 score :',f1)
print('The roc_auc_score is :',rc)
print('-----\n\nThe ROC curve :')
fig, _ = plt.subplots()
fpr,tpr,thresholds=roc_curve(y_test,y)
plt.plot([0, 1],[0, 1], '--')
plt.plot(fpr,tpr,label=kernel+' : '+str(rc))
plt.xlabel('false positive')
plt.ylabel('false negative')
plt.legend()
plt.show()
confusionMatrix = confusion_matrix(y_test, y)
print('-----\n\nThe confusion matrix is :')
ax = plt.axes()
sns.heatmap(confusionMatrix,annot=True)
ax.set_title('Confusion matrix of SVM '+str(kernel))
return yt_svm,yp_svm

# -----
# svm factor : factor affecting students performance, later on on this Ipython notebook we will explain how we will do this

# 1) factor as svm coefficients
def factors(array, K, max_or_min, df):

    n = array.shape[1]
    array = array.reshape(n,1)
    my_list = array.tolist()
    if max_or_min == 'max':
        temp = sorted(my_list)[-K:]
        res = []
        for ele in temp:

```

```

    n = array.shape[1]
    array = array.reshape(n,1)
    my_list = array.tolist()
    if max_or_min == 'max':
        temp = sorted(my_list)[-K:]
        res = []
        for ele in temp:
            res.append(my_list.index(ele))
        return(get_factors(res, df))

    elif max_or_min == 'min':
        temp = sorted(my_list, reverse=True)[-K:]
        temp = temp = np.array(temp).reshape(K,1)
        res = []
        for ele in temp:
            if ele<0:
                res.append(my_list.index(ele))
        return(get_factors(res, df))

    else:
        return

# 2) converts those factors to dataset columns name
def get_factors(index, df):
    f = []
    for i in index:
        f.append(df.columns[i])
    return f

```

```
# 3) Convert column names to understandable string
```

```
columns_name = {'famsize': 'family size', 'Pstatus': "parent's cohabitation status ", 'Medu': "mother's education",
                'Fedu': "father's education", 'Mjob': "mother's job", 'Fjob': "father's job",
                'reason': 'reason to choose this school ', 'schoolsup': 'extra educational support', 'famsup': 'family educational support',
                'paid': 'extra paid classes within the course subject', 'higher': 'wants to take higher education',
                'romantic': 'with a romantic relationship ', 'famrel': 'quality of family relationships', 'goout': 'going out with friends',
                'Dalc': 'workday alcohol consumption', 'Walc': 'weekend alcohol consumption'}
```

```
def column_to_string(fcts,max_or_min):
```

```
    if max_or_min == 'max':
        print('-----')
        print('Factors helping students succeed :')
    else:
        print('-----')
        print('-----')
        print('Factors leading students to failure')
```

```
    for fct in fcts:
        if fct in columns_name:
            print(columns_name[fct])
        else:
            print(fct)
```

```
def split(df,rest_size,test_size,randomState):
```

```
    data = df.to_numpy()
    n = data.shape[1]
    x = data[:,0:n-1]
    y = data[:,n-1]
    if(randomState):
        X_train,X_rest,y_train,y_rest = train_test_split(x,y,test_size=rest_size,random_state=randomState)
        X_val,X_test,y_val,y_test = train_test_split(X_rest,y_rest,test_size=test_size,random_state=randomState)
    else:
        X_train,X_rest,y_train,y_rest = train_test_split(x,y,test_size=rest_size,random_state=0)
        X_val,X_test,y_val,y_test = train_test_split(X_rest,y_rest,test_size=test_size,random_state=0)
```

```
    return X_train,X_val,X_test,y_train,y_val,y_test
# We will use the three different svm classifier kernels
# Linear kernel, polynomial kernel and gaussian kernel and we will choose the most accurate
```

```
##### Linear kernel #####
```

```
optimal_split_state1 = 0
maxAccuracy = 0
maxF1 = 0
```

```
# We already tune parameters, we do not need to loop over all the hyperparameters again,
# if you want to do so just set max_iteration to 2000 for example
# and remove the line 'optimal_split_state = 388628375' at the bottom of this cell.
```

```
max_iteration = 0
if max_iteration != 0:
    print ('-----Hyperparameters tuning starts-----\n\n')
```

```
for k in range(max_iteration):
    print ('Iteration :'+str(k)+', Current accuracy: '+str(maxAccuracy)+' Current f1 '+str(maxF1), end="\n")
    # let's get the optimal C value for the linear kernel
    split_state = np.random.randint(1,1000000000)-1
    X_train,X_val,X_test,y_train,y_val,y_test = split(df,rest_size=0.4,test_size=0.4,randomState=split_state)
    optimal_C = optimal_C_value()
```

```
# Now let's use the optimal C value
linear_clf = SVC(C=optimal_C,kernel='linear')
```

```
# Let's train the model with the optimal C value and calculate the training time
tic = time()
linear_clf.fit(X_train, y_train)
toc = time()
time1 = str(round(1000*(toc-tic))) + "ms"
y_linear = linear_clf.predict(X_test)
linear_f1 = f1_score(y_test, y_linear, average='macro')
linear_accuracy = accuracy_score(y_test, y_linear)*100
```

```

if linear_accuracy>maxAccuracy and linear_f1>maxF1:
    maxAccuracy = linear_accuracy
    maxF1 = linear_f1
    optimal_split_state1 = split_state
    if maxAccuracy>86 and maxF1>80:
        break;
# We've already tuned our hyperparameters, we will not repeat that again as it takes soo long.
# The optimal split state for linear kernel is 388628375
# Let's try that split state
optimal_split_state1 = 388628375
X_train,X_val,X_test,y_train,y_val,y_test = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state1)
optimal_C = optimal_C_value()

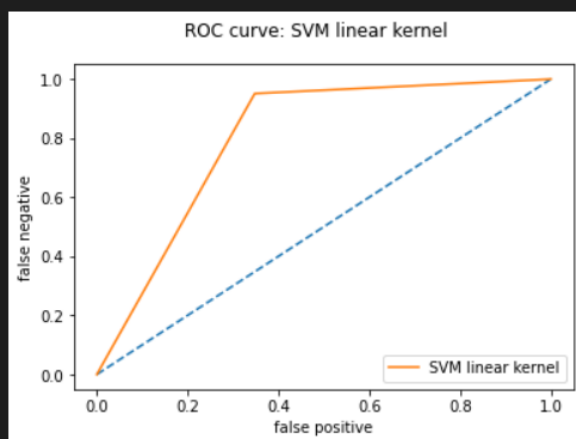
# Now let's use the optimal C value
linear_clf = SVC(C=optimal_C,kernel='linear')

# Let's train the model with the optimal C value and calculate the training time
tic = time()
linear_clf.fit(X_train, y_train)
toc = time()
time1 = str(round(1000*(toc-tic))) + "ms"
y_linear = linear_clf.predict(X_test)
linear_accuracy = accuracy_score(y_test, y_linear)*100
if max_iteration != 0:
    print('\n\n\n-----process ended'\n\n\n')

# Let's show the results
showResults(linear_accuracy, time1, y_linear,'SVM linear kernel')

```

The ROC curve is :



-----  
The model accuracy: 84 %

-----  
The training time is: 10ms


-----  
The f1 score is : 0.82

-----  
The roc\_auc\_score is : 0.8

-----



The confusion matrix is :



```
##### Polynomial kernel #####
optimal_split_state2 = 0
maxAccuracy = 0
maxF1 = 0

# We already tune parameters, we do not need to loop over all the hyperparameters again,
# if you want to do so just set max_iteration to 500 for example
# and remove the line 'optimal_split_state2 = 7070621' at the bottom of this cell.

max_iteration = 0
if max_iteration != 0:
    print ('-----Hyperparameters tuning starts-----\n\n')
for k in range(max_iteration):
    print ('Iteration :'+str(k)+', Current accuracy: '+str(maxAccuracy)+', Current f1 '+str(maxF1), end="\r")

    split_state = np.random.randint(1,100000000)-1
    X_train,X_val,X_test,y_train,y_val,y_test = split(df,rest_size=0.4,test_size=0.4,randomState=split_state)

    # Let's get the optimal C and the degree value for the polynomial kernel
    optimal_C, optimal_d = optimal_C_d_values()

    # Now let's use the optimal c value and the optimal degree value
    poly_clf = SVC(C=optimal_C,kernel='poly', degree=optimal_d)

    # Let's train the model with the optimal C value
    poly_clf.fit(X_train, y_train)
    y_poly = poly_clf.predict(X_test)
    poly_f1 = f1_score(y_test, y_poly, average='macro')
    poly_accuracy = accuracy_score(y_test, y_poly)*100
```

```
    if poly_accuracy>maxAccuracy and poly_f1>maxF1:
        maxAccuracy = poly_accuracy
        maxF1 = poly_f1
        optimal_split_state2 = split_state
# We've already tuned our hyperparameters, we will not repeat that again as it takes soo long.
# The optimal split state for polynomial kernel is 7070621
# Let's try that split state
optimal_split_state2 = 7070621

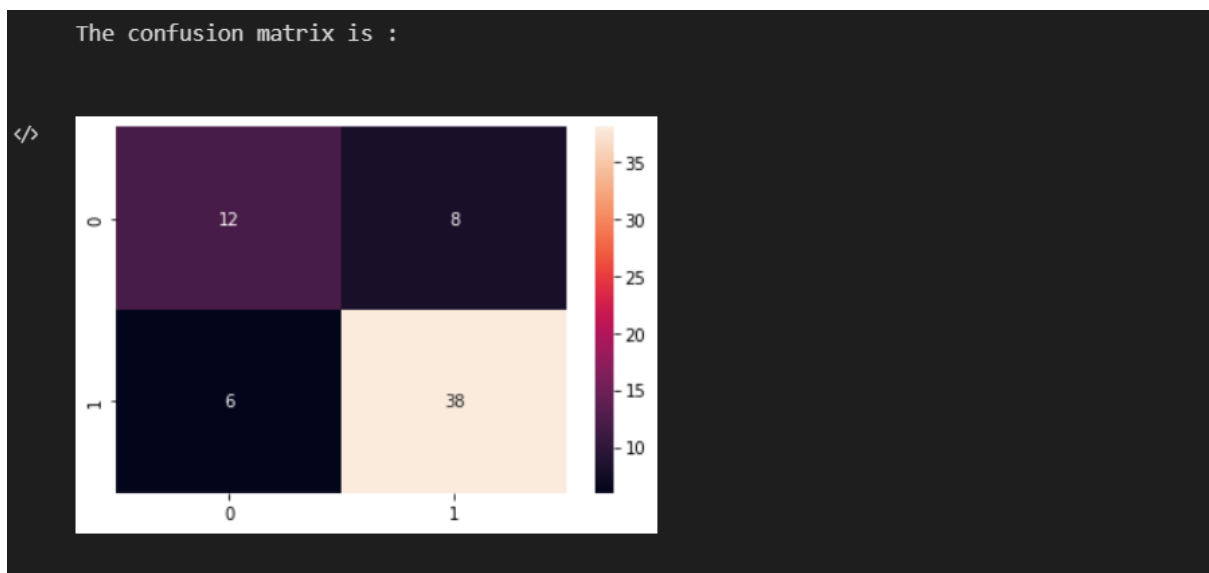

X_train,X_val,X_test,y_train,y_val,y_test = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state2)

optimal_C, optimal_d = optimal_C_d_values()

# Now let's use the optimal C value
poly_clf = SVC(C=optimal_C,kernel='poly', degree=optimal_d)

# Let's train the model and calculate the training time
tic = time()
poly_clf.fit(X_train, y_train)
toc = time()
time2 = str(round(1000*(toc-tic))) + "ms"
y_poly = poly_clf.predict(X_test)
poly_accuracy = accuracy_score(y_test, y_poly)*100
if max_iteration != 0:
    print('\n\n\n-----process ended'\n\n\n')

# Let's show the results
showResults(poly_accuracy, time2, y_poly,'SVM polynomial kernel')
```



```
##### Gaussian kernel #####
optimal_split_state3 = 0
maxAccuracy = 0
maxF1 = 0

# We already tune parameters, we do not need to loop over all the hyperparameters again,
# if you want to do so just set max_iteration to 500 for example
# and remove the line 'optimal_split_state3 = 93895097' at the bottom of this cell.

max_iteration = 0
if max_iteration != 0:
    print ('-----Hyperparameters tuning starts\'\\
    -----\\n\\n')
for k in range(max_iteration):
    print ('Iteration :'+str(k)+', Current accuracy: '+str(maxAccuracy)+', Current f1 '+str(maxF1), end="\\r")

    split_state = np.random.randint(1,100000000)-1
    x_train,x_val,x_test,y_train,y_val,y_test = split(df,rest_size=0.4,test_size=0.4,randomState=split_state)

    # Let's get the optimal C and the degree value for the polynomial kernel
    optimal_C, optimal_gamma = optimal_C_gamma_values()

    # Now let's use the optimal c value and the optimal degree value
    gauss_clf = SVC(C=optimal_C,kernel='rbf',gamma=optimal_gamma)

    # Let's train the model with the optimal C value
    gauss_clf.fit(X_train, y_train)
    y_gauss = gauss_clf.predict(X_test)
    gauss_f1 = f1_score(y_test, y_gauss, average='macro')
    gauss_accuracy = accuracy_score(y_test, y_gauss)*100
```

```

    if gauss_accuracy>maxAccuracy and gauss_f1>maxF1:
        maxAccuracy = gauss_accuracy
        maxF1 = gauss_f1
        optimal_split_state3 = split_state

# We've already tuned our hyperparameters, we will not repeat that again as it takes soo long.
# The optimal split state for polynomial kernel is 93895097
# Let's try that split state
optimal_split_state3 = 93895097
X_train,X_val,X_test,y_train,y_val,y_test = split(df,rest_size=0.4,test_size=0.4,randomState=optimal_split_state3)

optimal_C, optimal_gamma = optimal_C_gamma_values()

# Now let's use the optimal C value
gauss_clf = SVC(C=optimal_C,kernel='rbf',gamma=optimal_gamma)

# Let's train the model and calculate the training time
tic = time()
gauss_clf.fit(X_train, y_train)
toc = time()
time3 = str(round(1000*(toc-tic))) + "ms"
y_gauss = gauss_clf.predict(X_test)
gauss_accuracy = (accuracy_score(y_test, y_gauss)*100)


if max_iteration != 0:
    print('\n\n\n-----process ended'\n\n\n')

# Let's show the results
showResults(gauss_accuracy, time3, y_gauss,'SVM gaussian kernel')

```

... -----Results : SVM gaussian kernel -----

The ROC curve is :




-----  
The model accuracy: 83 %  
-----

-----  
The training time is: 5ms  
-----

-----  
The f1 score is : 0.77  
-----

The confusion matrix is :



```
compare_kernels()
```


```
[ ]
```


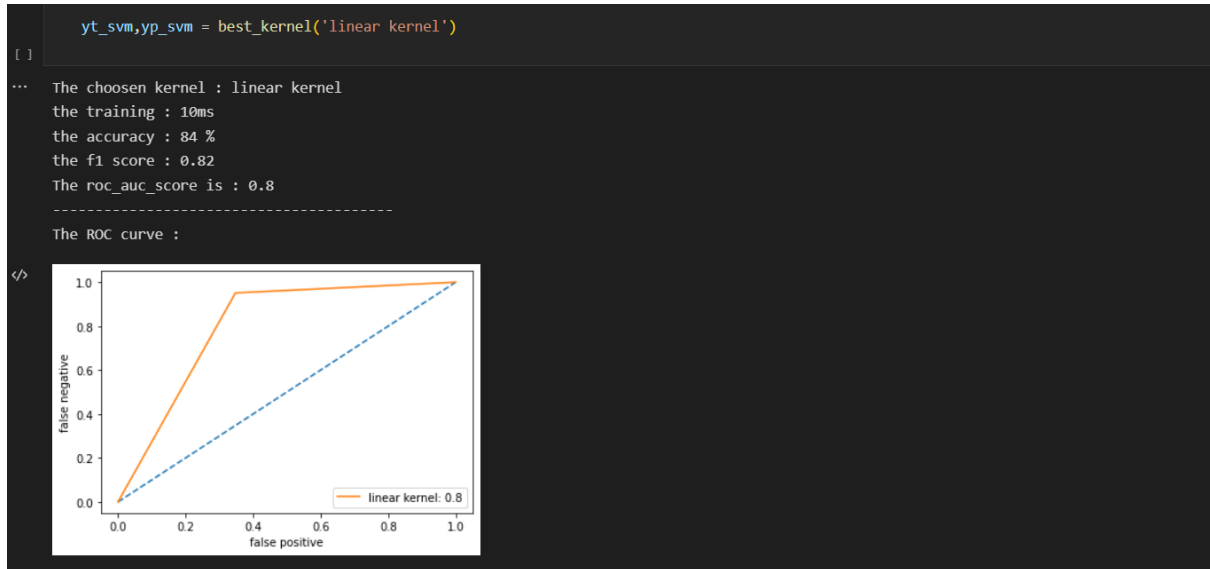
```
... ----- Comparison -----
```

metric	Linear kernel	polynomial kernel	gaussian kernel
training time	10ms	8ms	5ms
accuracy %	84.375	78.125	82.8125
confusion matrix	[15 8] [ 2 39]	[12 8] [ 6 38]	[10 11] [ 0 43]
f1 score	0.82	0.74	0.77
roc_auc_score	0.80	0.73	0.74

The Roc curves :

```
</>
```





```

# Get svm parameters
coefs = linear_clf.coef_

# factors helping students to succeed
column_to_string(factors(coefs, 5, 'max', df), 'max')

# factors leading students to failure
column_to_string(factors(coefs, 5, 'min', df), 'min')

```

[ ]

```

...
-----
Factors helping students succeed :
father's education
guardian
wants to take higher education
studytime
father's job
-----
-----

```

```

Factors leading students to failure
age
health
going out with friends
absences
failures

```

# Function to compare the three classifiers (Logistic regression, KNN and SVM) performances :

```

def compare_lg_knn_svm(yt_knn, yp_knn, yt_lg, yp_lg, yt_svm, yp_svm):
    #F1 score
    f1_lg = round(f1_score(yt_lg, yp_lg, average='macro')*100)
    f1_knn = round(f1_score(yt_knn, yp_knn, average='macro')*100)
    f1_svm = round(f1_score(yt_svm, yp_svm, average='macro')*100)

    #Accuracy score
    acc_lg = round(accuracy_score(yt_lg, yp_lg)*100)
    acc_knn = round(accuracy_score(yt_knn, yp_knn)*100)
    acc_svm = round(accuracy_score(yt_svm, yp_svm)*100)

    #Confusion matrix
    conf_lg = confusion_matrix(yt_lg, yp_lg)
    conf_knn = confusion_matrix(yt_knn, yp_knn)
    conf_svm = confusion_matrix(yt_svm, yp_svm)

    #ROC score
    roc_c_lg = round(roc_auc_score(yt_lg, yp_lg)*100)
    roc_c_knn = round(roc_auc_score(yt_knn, yp_knn)*100)
    roc_c_svm = round(roc_auc_score(yt_svm, yp_svm)*100)

    #ROC curve thresholds
    roc_knn = roc_curve(yt_knn, yp_knn)
    roc_lg = roc_curve(yt_lg, yp_lg)
    roc_svm = roc_curve(yt_svm, yp_svm)

```

```

# Table of metrics
print('-----Table of metrics-----\n\n')
data_rows = [('f1 score',f1_lg,f1_knn,f1_svm),
              ('', '', '', ''),
              ('accuracy %',acc_lg,acc_knn,acc_svm),
              ('', '', '', ''),
              ('confusion matrix',conf_lg[0], conf_knn[0], conf_svm[0]),
              ('',conf_lg[1], conf_knn[1], conf_svm[1]),
              ('', '', '', ''),
              ('ROC score',roc_c_lg,roc_c_knn,roc_c_svm)]
t = Table(rows=data_rows, names=('metric','Logistic regression', 'KNN', 'SVM'))
print(t)

#Plot ROC curve
print('\n\n-----ROC curves-----\n\n')
fig, _ = plt.subplots()
fig.suptitle('Comparison of three ROC curves')
fpr,tpr,thresholds=roc_lg
plt.plot([0, 1],[0, 1], '--')
plt.plot(fpr,tpr,label='Logistic regression :'+str(roc_c_lg))
plt.xlabel('false positive')
plt.ylabel('false negative')
fpr,tpr,thresholds=roc_knn
plt.plot(fpr,tpr,label='KNN :'+str(roc_c_knn))
fpr,tpr,thresholds=roc_svm
plt.plot(fpr,tpr,label='SVM :'+str(roc_c_svm))
plt.legend()
plt.show()

```

```

# Maximum metrics
print('-----Max of metrics-----\n\n')
data_rows = [('max f1 score',algo_with_max_metric(f1_lg,f1_knn,f1_svm)),
              ('', '', '', ''),
              ('max accuracy %',algo_with_max_metric(acc_lg,acc_knn,acc_svm)),
              ('', '', '', ''),
              ('max ROC score',algo_with_max_metric(roc_c_lg,roc_c_knn,roc_c_svm))]
t = Table(rows=data_rows, names=('metric','Learning algorithm winnig'))
print(t)

# Function returning name of winnig algorithm based on a single metric
def algo_with_max_metric(a,b,c):
    max_metric = max(a,b,c)
    if max_metric == a:
        return 'Logistic regression'
    elif max_metric == b:
        return 'KNN'
    else:
        return 'SVM'


```

```
compare_lg_knn_svm(yt_knn,yp_knn,yt_lg,yp_lg,yt_svm,yp_svm)
```

```
''' -----Table of metrics-----
```

metric	Logistic regression	KNN	SVM
f1 score	74	48	82
accuracy %	81	70	84
confusion matrix	[18 17] [ 6 78]	[ 3 23] [13 80]	[15 8] [ 2 39]
ROC score	72	49	80

-----ROC curves-----



-----Max of metrics-----

metric	Learning algorithm winnig
max f1 score	SVM
max accuracy %	SVM
max ROC score	SVM

### INFERENCE:

- Here we have used three models such as logistic regression, KNN and SVM to check the performance of those models and identify which model performs well with better accuracy.
- The models are hyper tuned to increase their accuracy.
- From this it is identified that SVM provides greater accuracy compared to all other models
- In SVM we identified the accuracy of three different kernels such as linear, polynomial and gaussian kernel among the three the linear kernel of SVM provides greater accuracy of 84%